

Efficient Stop & Warning Sign and Pedestrian Detection

Castleberry, Cherry, and Firth

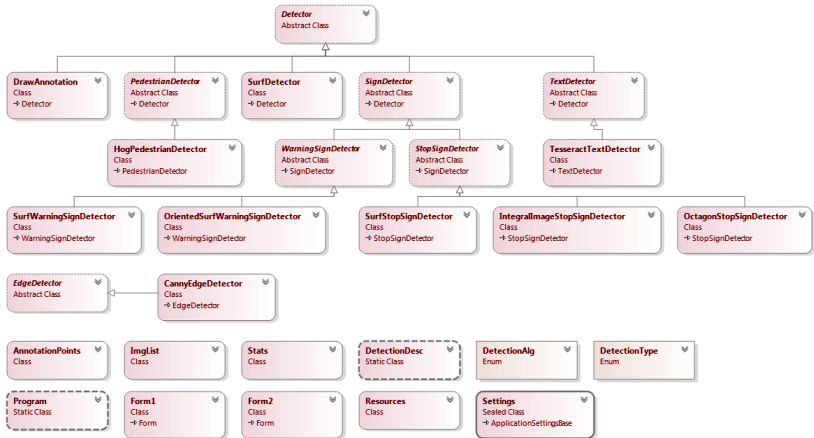
December 6, 2012

Overview

- We implemented the following:
 - Stop-sign detector, using the notion of integral images from SURF
 - Warning sign detector, using perspective transformation on a model image, then SURF

Wrappers

- We created a wrapper to the pedestrian detector to gain familiarity with the EmguCV library.
- We also created a wrapper to the built-in EmguCV SURF detector to use as a basis of comparison for our own methods.







Stop Sign Detection, Integral Images

- We developed a method for detecting stop signs based upon the use of integral images which we encountered in the SURF algorithm.
- To summarize the method, we use integral images from both the left-hand side (top left) and the right-hand side (bottom right) on only the R channel. Then we consider only the LHS and RHS integral images along the diagonal of the image. We difference these, then fit a Gaussian curve to the resulting vector. We threshold the curve at the points of inflection to generate a bounding box for the sign.




Stop Sign Detection, Integral Images

- To begin our method, we scale the $N \times M$ image to $N \times N$ for $N < M$, $M \times M$ for $M < N$. We require a square matrix to extract a particular vector.
- Recall that the formula for computing the integral image \mathcal{I}_- at a pixel (x, y) with intensity value $I(x, y)$ is:

$$\mathcal{I}_-(x, y) = \sum_{i=0}^{n_x} \sum_{j=0}^{n_y} I(x, y) \quad (1)$$

- We compute the integral image at a pixel using the following formula:



$$\mathcal{I}_{x,y} = \mathcal{I}_{x-1,y} + \mathcal{I}_{x,y-1} - \mathcal{I}_{x-1,y-1} \quad (2)$$



Stop Sign Detection, Integral Images

- Likewise, we compute an RHS integral image \mathcal{I}_+ at a pixel (x, y) with intensity value $I(x, y)$ as:

$$\mathcal{I}_+(x, y) = \sum_{i=N}^{n_x} \sum_{j=N}^{n_y} I(x, y) \quad (3)$$



Stop Sign Detection, Integral Images


- After obtaining the integral image, we copy its diagonal into a vector u .
- We then apply a finite-difference method to the elements in u and store it in v , as follows:


$$v_n = u_n - u_{n-1} \quad (4)$$

The vector v gives the LHS crosshair of the R-channel. For images which have stop signs, v has a Gauss distribution.

Stop Sign Detection, Integral Images

- We apply this finite-difference method for both vectors u_- and u_+ to obtain v_- and v_+ .
- Then, we add v_- and v_+ to obtain a vector m .
- Finally, we compute the standard deviation σ of the vector m and its centroid c , then apply a Gaussian fit to the data in m . We call the Gaussian fit G .
- We then apply finite-differencing to the Gaussian fit G to obtain G' , then find the indices at $\min(G')$ and $\max(G')$.
- These indices form the bounding box for the image.





Warning Signs

- For our warning sign detection, we experimented with SURF in combination with perspective transformations on out-of-plane-rotated signs.
- In particular, we assume that we know the perspective information of an out-of-plane-rotated sign. We apply a perspective transform to a model sign, then apply SURF to the two images, then match.



Warning Signs


- We first hand-annotated N images using a MATLAB code.
- From this, we extracted a set of $4N$ points which give the vertices of the warning sign. We then applied `getPerspectiveTransform()` on the points, which returns a perspective transform matrix M .
- We apply this perspective transform matrix to the model with the function `warpPerspective()`.
- We then run SURF on the two images, then compare their feature descriptors to obtain a match within a threshold of t_{ϵ} .

Perspective Transform

- If $a_{x,y,z}$ is the point to be projected, $c_{x,y,z}$ is the camera, $\theta_{x,y,z}$ is the camera orientation and $e_{x,y,z}$ is the position of the viewer relative to the display surface then $b_{x,y}$, the 2D projection of a , is given by:

$$\begin{bmatrix} d_x \\ d_y \\ d_z \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 \\ 0 & \cos(\theta_x) & -\sin(\theta_x) \\ 0 & \sin(\theta_x) & \cos(\theta_x) \end{bmatrix} \begin{bmatrix} \cos(\theta_y) & 0 & \sin(\theta_y) \\ 0 & 1 & 0 \\ -\sin(\theta_y) & 0 & \cos(\theta_y) \end{bmatrix} \begin{bmatrix} \cos(\theta_z) & -\sin(\theta_z) & 0 \\ \sin(\theta_z) & \cos(\theta_z) & 0 \\ 0 & 0 & 1 \end{bmatrix} \left(\begin{bmatrix} a_x \\ a_y \\ a_z \end{bmatrix} - \begin{bmatrix} c_x \\ c_y \\ c_z \end{bmatrix} \right). \quad (5)$$

- A visual example follows.



Results

- We considered there to be a match if at least 50% of the sign fell within the area of the bounding box.
- The following table summarizes our results:

Surf Stop Sign Detector	17%
Integral Images Stop Sign Detector	67%
Surf Warning Sign Detector	42%
Oriented Surf Warning Sign Detector	62%

- As is evident in the above, our accuracy for . . . was an abysmal $N\%$.

Conclusion

- Bearing in mind that the accuracy for our integral image detector was much less than the accuracy for the built-in SURF algorithm, we conclude that in the development of computer vision algorithms, one should:
 - 1 Not over-utilize heuristics, and
 - 2 Develop the algorithm starting from optimized versions of existing algorithms.